

```
from google.colab import files
uploaded = files.upload()
import pandas as pd
import io
```

Choose files UCI Heart Disease.csv
UCI Heart Disease Data.csv(text/csv) - 79346 bytes, last modified: 10/12/2025 - 100% done
 Saving UCI Heart Disease Data.csv to UCI Heart Disease Data (1).csv

```
df = pd.read_csv(io.BytesIO(uploaded["UCI Heart Disease Data (1).csv"]))
print("Loaded:", df.shape)
```

Loaded: (920, 16)

```
import pandas as pd
df = pd.read_csv(list(uploaded.keys())[0])
df.head()
```

	id	age	sex	dataset	cp	trestbps	chol	fb	restecg	thalch	exang	oldpeak	slope	ca
0	1	63	Male	Cleveland	typical angina	145.0	233.0	True	lv hypertrophy	150.0	False	2.3	downsloping	0.0
1	2	67	Male	Cleveland	asymptomatic	160.0	286.0	False	lv hypertrophy	108.0	True	1.5	flat	3.0
2	3	67	Male	Cleveland	asymptomatic	120.0	229.0	False	lv hypertrophy	129.0	True	2.6	flat	2.0

Next steps: [Generate code with df](#) [New interactive sheet](#)

Importing the Dependencies

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn.metrics import classification_report, confusion_matrix, roc_curve, auc
from sklearn.model_selection import learning_curve
```

A. Data Cleaning

```
df = df.drop_duplicates()
df = df.replace(["", " ", "NA", "NaN"], np.nan)
```

```
num_cols = ["age", "trestbps", "chol", "thalch", "oldpeak", "ca", "num"]
for col in num_cols:
    if col in df.columns:
        df[col] = pd.to_numeric(df[col], errors="coerce")
```

```
if "chol" in df.columns:
    df.loc[df["chol"] == 0, "chol"] = np.nan
```

```
for col in df.columns:
    if pd.api.types.is_numeric_dtype(df[col]):
        df[col] = df[col].fillna(df[col].median())
    else:
        mode = df[col].mode()
        if not mode.empty:
            df[col] = df[col].fillna(mode.iloc[0])
```

```
/tmp/ipython-input-2363592276.py:7: FutureWarning: Downcasting object dtype arrays on .fillna, .ffill, .bfill is
df[col] = df[col].fillna(mode.iloc[0])
```

B. Encoding

```

if "sex" in df.columns:
    df["sex"] = df["sex"].astype(str).str.upper().map({"MALE":1,"FEMALE":0}).fillna(0).astype(int)
for c in ["fbs","exang"]:
    if c in df.columns:
        df[c] = df[c].astype(str).str.upper().map({"TRUE":1,"FALSE":0}).fillna(0).astype(int)

if "thal" in df.columns:
    df["thal"] =
        df["thal"].astype(str).str.strip()
        .replace({"reversible defect":"reversible defect"})
    )

cat_targets = ["cp","restecg","slope","thal","dataset"]
cat_cols = [c for c in cat_targets if c in df.columns]
for c in cat_cols:
    df[c] = df[c].astype(str).str.strip()
df = pd.get_dummies(df, columns=cat_cols, drop_first=True)

```

C.Feature Scaling

```

y = (df["num"] > 0).astype(int)
X = df.drop(columns=[c for c in ["id","num"] if c in df.columns])

X = X.replace([np.inf, -np.inf], np.nan).fillna(0)

```

D. Train- Test Split

```

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

print("Preprocessing OK")
print("Train:", X_train_scaled.shape, "| Test:", X_test_scaled.shape, "| Features:", X.shape[1])

Preprocessing OK
Train: (736, 21) | Test: (184, 21) | Features: 21

```

Apply 5 Machine Learning Algo.

```

# Logistic Regression
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, confusion_matrix
log_model = LogisticRegression(max_iter=1000)
log_model.fit(X_train_scaled, y_train)
y_pred_log = log_model.predict(X_test_scaled)
print("\n Logistic Regression Results")
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred_log))
print("\nClassification Report:\n", classification_report(y_test, y_pred_log, digits=4))

```

```

Logistic Regression Results
Confusion Matrix:
[[65 17]
 [10 92]]

Classification Report:
      precision    recall   f1-score   support
          0       0.8667    0.7927    0.8280      82
          1       0.8440    0.9020    0.8720     102

      accuracy                           0.8533      184
     macro avg       0.8554    0.8473    0.8500      184
  weighted avg       0.8541    0.8533    0.8524      184

```

K-Nearest Neighbors

```

from sklearn.neighbors import KNeighborsClassifier
knn_model = KNeighborsClassifier(n_neighbors=7)

```

```
knn_model.fit(X_train_scaled, y_train)
y_pred_knn = knn_model.predict(X_test_scaled)

print("\nQ KNN Results")
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred_knn))
print("\nClassification Report:\n", classification_report(y_test, y_pred_knn, digits=4))
```

Q KNN Results
 Confusion Matrix:
 [[67 15]
 [9 93]]

Classification Report:				
	precision	recall	f1-score	support
0	0.8816	0.8171	0.8481	82
1	0.8611	0.9118	0.8857	102
accuracy			0.8696	184
macro avg	0.8713	0.8644	0.8669	184
weighted avg	0.8702	0.8696	0.8690	184

```
# Decision Tree

from sklearn.tree import DecisionTreeClassifier
dt_model = DecisionTreeClassifier(max_depth=5, random_state=42)
dt_model.fit(X_train_scaled, y_train)
y_pred_dt = dt_model.predict(X_test_scaled)
print("\nQ Decision Tree Results")
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred_dt))
print("\nClassification Report:\n", classification_report(y_test, y_pred_dt, digits=4))
```

Q Decision Tree Results
 Confusion Matrix:
 [[69 13]
 [20 82]]

Classification Report:				
	precision	recall	f1-score	support
0	0.7753	0.8415	0.8070	82
1	0.8632	0.8039	0.8325	102
accuracy			0.8207	184
macro avg	0.8192	0.8227	0.8198	184
weighted avg	0.8240	0.8207	0.8211	184

```
# Random Forest

from sklearn.ensemble import RandomForestClassifier
rf_model = RandomForestClassifier(n_estimators=200, random_state=42)
rf_model.fit(X_train_scaled, y_train)
y_pred_rf = rf_model.predict(X_test_scaled)
print("\nQ Random Forest Results")
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred_rf))
print("\nClassification Report:\n", classification_report(y_test, y_pred_rf, digits=4))
```

Q Random Forest Results
 Confusion Matrix:
 [[68 14]
 [15 87]]

Classification Report:				
	precision	recall	f1-score	support
0	0.8193	0.8293	0.8242	82
1	0.8614	0.8529	0.8571	102
accuracy			0.8424	184
macro avg	0.8403	0.8411	0.8407	184
weighted avg	0.8426	0.8424	0.8425	184

```
# Support Vector Machine

from sklearn.svm import SVC
svm_model = SVC(kernel="rbf", probability=True, random_state=42)
```

```

svm_model.fit(X_train_scaled, y_train)
y_pred_svm = svm_model.predict(X_test_scaled)
print("\nSVM Results")
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred_svm))
print("\nClassification Report:\n", classification_report(y_test, y_pred_svm, digits=4))

```

```

SVM Results
Confusion Matrix:
[[63 19]
 [ 6 96]]

Classification Report:
precision    recall   f1-score   support
          0    0.9130    0.7683    0.8344      82
          1    0.8348    0.9412    0.8848     102

accuracy                           0.8641      184
macro avg                           0.8739    0.8547    0.8596      184
weighted avg                          0.8697    0.8641    0.8624      184

```

Model result Analysis

```

# Logistic Regression Evaluation

from sklearn.metrics import confusion_matrix, classification_report, roc_curve, auc
from sklearn.model_selection import learning_curve
import matplotlib.pyplot as plt
import numpy as np
y_pred_log = log_model.predict(X_test_scaled)

#Confusion Matrix
print("Confusion Matrix (Logistic Regression):\n", confusion_matrix(y_test, y_pred_log))

#Precision, Recall, F1-score
print("\nClassification Report:\n", classification_report(y_test, y_pred_log, digits=4))

#ROC Curve
y_prob_log = log_model.predict_proba(X_test_scaled)[:,1]
fpr, tpr, _ = roc_curve(y_test, y_prob_log)
plt.plot(fpr, tpr, label=f"Logistic (AUC={auc(fpr,tpr):.3f})")
plt.plot([0,1],[0,1],"k--")
plt.title("ROC Curve - Logistic Regression")
plt.xlabel("False Positive Rate"); plt.ylabel("True Positive Rate")
plt.legend(); plt.show()



#Learning Curve
train_sizes, train_scores, val_scores = learning_curve(
    log_model, X_train_scaled, y_train, cv=5,
    train_sizes=np.linspace(0.1,1.0,5), scoring="accuracy"
)
plt.plot(train_sizes, train_scores.mean(axis=1), "o-", label="Training")
plt.plot(train_sizes, val_scores.mean(axis=1), "o-", label="Validation")
plt.title("Learning Curve - Logistic Regression")
plt.xlabel("Training samples"); plt.ylabel("Accuracy")
plt.legend(); plt.show()

```

```
Confusion Matrix (Logistic Regression):
[[65 17]
 [10 92]]
```

```
Classification Report:
precision    recall   f1-score   support
          0    0.8667   0.7927   0.8280      82
          1    0.8440   0.9020   0.8720     102

   accuracy          0.8554   0.8473   0.8500     184
macro avg    0.8554   0.8473   0.8500     184
weighted avg  0.8541   0.8533   0.8524     184
```



```
# KNN Evaluation
y_pred_knn = knn_model.predict(X_test_scaled)

print("Confusion Matrix (KNN):\n", confusion_matrix(y_test, y_pred_knn))
print("\nClassification Report:\n", classification_report(y_test, y_pred_knn, digits=4))

#ROC Curve
y_prob_knn = knn_model.predict_proba(X_test_scaled)[:,1]
fpr, tpr, _ = roc_curve(y_test, y_prob_knn)
plt.plot(fpr, tpr, label=f"KNN (AUC={auc(fpr,tpr):.3f})")
plt.plot([0,1],[0,1],"k--")
plt.title("ROC Curve - KNN"); plt.xlabel("FPR"); plt.ylabel("TPR")
plt.legend(); plt.show()



#Learning Curve
train_sizes, train_scores, val_scores = learning_curve(
    knn_model, X_train_scaled, y_train, cv=5,
    train_sizes=np.linspace(0.1,1.0,5), scoring="accuracy"
)
```

```
plt.plot(train_sizes, train_scores.mean(axis=1), "o-", label="Training")
plt.plot(train_sizes, val_scores.mean(axis=1), "o-", label="Validation")
plt.title("Learning Curve - KNN")
plt.xlabel("Training samples"); plt.ylabel("Accuracy")
plt.legend(); plt.show()
```

Confusion Matrix (KNN):
[[67 15]
 [9 93]]

Classification Report:

	precision	recall	f1-score	support
0	0.8816	0.8171	0.8481	82
1	0.8611	0.9118	0.8857	102
accuracy			0.8696	184
macro avg	0.8713	0.8644	0.8669	184
weighted avg	0.8702	0.8696	0.8690	184



```
# Decision Tree Evaluation

y_pred_dt = dt_model.predict(X_test_scaled)

print("Confusion Matrix (Decision Tree):\n", confusion_matrix(y_test, y_pred_dt))
print("\nClassification Report:\n", classification_report(y_test, y_pred_dt, digits=4))

#ROC Curve
y_prob_dt = dt_model.predict_proba(X_test_scaled)[:,1]
fpr, tpr, _ = roc_curve(y_test, y_prob_dt)
plt.plot(fpr, tpr, label=f"Decision Tree (AUC={auc(fpr,tpr):.3f})")
plt.plot([0,1],[0,1],"k--")
plt.title("ROC Curve - Decision Tree"); plt.xlabel("FPR"); plt.ylabel("TPR")
```

```

plt.legend(); plt.show()

#Learning Curve
train_sizes, train_scores, val_scores = learning_curve(
    dt_model, X_train_scaled, y_train, cv=5,
    train_sizes=np.linspace(0.1,1.0,5), scoring="accuracy"
)
plt.plot(train_sizes, train_scores.mean(axis=1), "o-", label="Training")
plt.plot(train_sizes, val_scores.mean(axis=1), "o-", label="Validation")
plt.title("Learning Curve - Decision Tree")
plt.xlabel("Training samples"); plt.ylabel("Accuracy")
plt.legend(); plt.show()

```

Confusion Matrix (Decision Tree):

```


[[69 13]
 [20 82]]

```


Classification Report:

	precision	recall	f1-score	support
0	0.7753	0.8415	0.8070	82
1	0.8632	0.8039	0.8325	102
accuracy			0.8207	184
macro avg	0.8192	0.8227	0.8198	184
weighted avg	0.8240	0.8207	0.8211	184

ROC Curve - Decision Tree



Learning Curve - Decision Tree



```
# Random Forest Evaluation
```

```

y_pred_rf = rf_model.predict(X_test_scaled)

print("Confusion Matrix (Random Forest):\n", confusion_matrix(y_test, y_pred_rf))
print("\nClassification Report:\n", classification_report(y_test, y_pred_rf, digits=4))

```



```
#ROC Curve
y_prob_rf = rf_model.predict_proba(X_test_scaled)[:,1]
fpr, tpr, _ = roc_curve(y_test, y_prob_rf)
plt.plot(fpr, tpr, label=f"Random Forest (AUC={auc(fpr,tpr):.3f})")
plt.plot([0,1],[0,1],"k--")
plt.title("ROC Curve - Random Forest"); plt.xlabel("FPR"); plt.ylabel("TPR")
plt.legend(); plt.show()

#Learning Curve
train_sizes, train_scores, val_scores = learning_curve(
    rf_model, X_train_scaled, y_train, cv=5,
    train_sizes=np.linspace(0.1,1.0,5), scoring="accuracy"
)
plt.plot(train_sizes, train_scores.mean(axis=1), "o-", label="Training")
plt.plot(train_sizes, val_scores.mean(axis=1), "o-", label="Validation")
plt.title("Learning Curve - Random Forest")
plt.xlabel("Training samples"); plt.ylabel("Accuracy")
plt.legend(); plt.show()
```

Confusion Matrix (Random Forest):

```
[[68 14]
 [15 87]]
```

	precision	recall	f1-score	support
0	0.8193	0.8293	0.8242	82
1	0.8614	0.8529	0.8571	102
accuracy			0.8424	184
macro avg	0.8403	0.8411	0.8407	184
weighted avg	0.8426	0.8424	0.8425	184



```
# SVM Evaluation

y_pred_svm = svm_model.predict(X_test_scaled)

print("Confusion Matrix (SVM):\n", confusion_matrix(y_test, y_pred_svm))
print("\nClassification Report:\n", classification_report(y_test, y_pred_svm, digits=4))



#ROC Curve
y_prob_svm = svm_model.predict_proba(X_test_scaled)[:,1]
fpr, tpr, _ = roc_curve(y_test, y_prob_svm)
plt.plot(fpr, tpr, label=f"SVM (AUC={auc(fpr,tpr):.3f})")
plt.plot([0,1],[0,1],"k--")
plt.title("ROC Curve - SVM"); plt.xlabel("FPR"); plt.ylabel("TPR")
plt.legend(); plt.show()

#Learning Curve
train_sizes, train_scores, val_scores = learning_curve(
    svm_model, X_train_scaled, y_train, cv=5,
    train_sizes=np.linspace(0.1,1.0,5), scoring="accuracy"
)
plt.plot(train_sizes, train_scores.mean(axis=1), "o-", label="Training")
plt.plot(train_sizes, val_scores.mean(axis=1), "o-", label="Validation")
plt.title("Learning Curve - SVM")
plt.xlabel("Training samples"); plt.ylabel("Accuracy")
plt.legend(); plt.show()
```

```
Confusion Matrix (SVM):
[[63 19]
 [ 6 96]]
```


```
Classification Report:
precision    recall   f1-score   support
          0    0.9130    0.7683    0.8344      82
          1    0.8348    0.9412    0.8848     102

   accuracy                           0.8641    184
  macro avg    0.8739    0.8547    0.8596    184
weighted avg   0.8697    0.8641    0.8624    184
```




Visualization Requirements

```
#Histograms
fig, axes = plt.subplots(1, 3, figsize=(16, 4))
for ax, col in zip(axes, ["age", "trestbps", "chol"]):
    if col in df.columns:
        sns.histplot(df[col], bins=30, kde=True, ax=ax, color="skyblue")
        ax.set_title(f"{col} Distribution")
plt.tight_layout()
plt.show()
```



```
#Correlation Heatmap
corr = df.select_dtypes(include=[np.number]).corr()
plt.figure(figsize=(12, 10))
sns.heatmap(corr, cmap="coolwarm", center=0, annot=False)
plt.title("Correlation Heatmap")
plt.show()
```



```
#Scatter Plots
if {"age","chol"}.issubset(df.columns):
    plt.figure(figsize=(6, 5))
    sns.scatterplot(x=df["age"], y=df["chol"], hue=y, palette="Set1", alpha=0.8)
    plt.title("Age vs Cholesterol (colored by disease)")
    plt.show()

if {"trestbps","thalch"}.issubset(df.columns):
    plt.figure(figsize=(6, 5))
    sns.scatterplot(x=df["trestbps"], y=df["thalch"], hue=y, palette="Set2", alpha=0.8)
    plt.title("Resting BP vs Max Heart Rate (colored by disease)")
    plt.show()
```

